

```
In [1]: import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import datasets, transforms
from torch.utils.data import DataLoader, random_split
import numpy as np
import time
```

## Introduction

This file outlines the steps used in training a Pytorch model on EMNIST letter dataset. Test results are stored in Pytoch\_Testing.ipynb

## Load Dataset

Define class labels A~Z, the first member here is meaningless

```
In [2]: classes = []
classes.append(chr(48))
for i in range(65, 91):
    classes.append(chr(i))
class_size = len(classes)

# 26 letters
display(np.transpose(classes))
print("class size = ", class_size)

array(['0', 'A', 'B', 'C', 'D', 'E', 'F', 'G', 'H', 'I', 'J', 'K', 'L',
       'M', 'N', 'O', 'P', 'Q', 'R', 'S', 'T', 'U', 'V', 'W', 'X', 'Y',
       'Z'], dtype='<U1')
class size = 27
```

```
In [3]: import pandas as pd
from pandas.plotting import table
import matplotlib.pyplot as plt

pd.set_option('display.max_columns', None)
pd.set_option('display.expand_frame_repr', False)
pd.set_option('max_colwidth', -1)
```

```
In [4]: df = pd.DataFrame(classes, columns=['classes'])
display(df.T)
```

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26
classes	0	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W	X	Y	Z

## Construct custom transform

Define transforms to dataset:

1. Convert to tensor model
2. Normalize data

```
In [5]: transform = transforms.Compose(
        [transforms.ToTensor(),
         transforms.Normalize((0.5), (0.5))])
```

## Load data

When running the following command for the first time, pytorch will prompt you to download the EMNIST dataset. Other available datasets are byclass, bymerge, balanced, letters, digits, mnist. Use different dataset by setting `split`.

```
In [6]: batch_size = class_size #load one batch per group(A~Z)

# training set
trainset = datasets.EMNIST('data', download=True, \
                           train=True, transform=transform, split='letters')
trainloader = DataLoader(trainset, batch_size=batch_size, shuffle=True, num_workers=2)

testset = datasets.EMNIST('data', download=True, \
                          train=False, transform=transform, split='letters')

# separate testset into 50% for testing and 50% for validation
testset, validset = random_split(testset, [int(len(testset)/2), int(len(testset)/2)])
testloader = DataLoader(testset, batch_size=batch_size, shuffle=True, num_workers=2)
validloader = DataLoader(validset, batch_size=batch_size, shuffle=True, num_workers=2)
```

## Display and check data

The size of each image in data is 28x28.

Also confirm that 26 images are grabbed from one batch.

```
In [7]: dataiter = iter(trainloader)
images, labels = dataiter.next()

print(images.shape)
```

```
torch.Size([27, 1, 28, 28])
```

Below provides two ways of displaying one batch of dataset.

1. `imshow` displays one batch without label, so remember to print out labels after calling this function.
2. `display_batch` displays images with corresponding labels as title, but it processes images slower than `imshow`


```
In [8]: import matplotlib.pyplot as plt
from torchvision.utils import make_grid

def imshow(img):
    """ Helper function for displaying one batch of images without labels
    Notice that due to the flipping and rotating when processing EMNIST images, the order
    of grid is from up to down then from left to right.

    Args:
        img: one batch of images
    """
    img = make_grid(images)
    img = img / 2 + 0.5
    npimg = img.numpy()

    # EMNIST dataset is mirrored and rotated
    # convert the images back to regular view
    npimg = np.fliplr(npimg)
    npimg = np.rot90(npimg, axes=(-1,-2))
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()
```

```
In [9]: imshow(images)
print("Ground Truth: ", ' '.join('%5s' % classes[labels[j]] for j in range(batch_size)))
```



Ground Truth: T K Y X C I M Q A N V Q E C T  
 X N A F U O I I X Y T M

```
In [10]: def display_batch(img, labels, n_row=5, n_col=6):
  """ Helper function for displaying one batch of images.
  The default batch size is 26.


  Args:
      img: one batch of images
      labels: ground truth or predicted label of images
      n_row: number of rows of subplots
      n_col: number of columns of subplots
  """
  fig, ax = plt.subplots(n_row, n_col)
  fig.tight_layout(pad=0.1)
  row = 0
  for i, img in enumerate(images):
      if i % n_col == 0 and i > 0:
          row += 1

      # EMNIST dataset is mirrored and rotated
      # convert the images back to regular view
      npimg = img.numpy()
      npimg = np.fliplr(npimg)
      npimg = np.rot90(npimg, axes=(-1, -2))
      npimg = np.transpose(npimg, (1, 2, 0))

      # configure into subplots
      sub_plot = ax[row, i % n_col]
      sub_plot.imshow(npimg)
      sub_plot.title.set_text(classes[labels[i]])
      # disable xy axis to make room for labels
      sub_plot.get_xaxis().set_visible(False)
      sub_plot.get_yaxis().set_visible(False)
  for i in range(n_row * n_col - len(labels), n_row * n_col):
      # remove unused blocks
      sub_plot = ax[row, i % n_col]
      sub_plot.set_axis_off()

  plt.show()
```

```
In [11]: display_batch(images, labels)
```



## Define Neural Network

```
In [12]: class Net(nn.Module):
    """ Neural network for training and classifying EMNIST dataset.

    Network architecture:
    - max pooling
    - 2D convolution layer of 64 channels and kernel size 5
    - 2D convolution layer of 128 channels and kernel size 5
    - Input layer
    - First hidden layer: fully connected layer of size 128 nodes
    - Second hidden layer: fully connected layer of size 64 nodes
    - Output layer: a linear layer with one node per class

    Activation function: ReLU for all layers
    """
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 64, kernel_size=5)
        self.conv2 = nn.Conv2d(64, 128, kernel_size=5)
        self.fc1 = nn.Linear(128*4*4, 128)
        self.fc2 = nn.Linear(128, 64)
        self.fc3 = nn.Linear(64, class_size)

    def forward(self, x):
        self.pool = nn.MaxPool2d(2, 2)
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 128*4*4)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()
```

## Define Loss Function and Optimizer

Use default mean for cross-entropy.

Learning rate is 0.001.

Momentum is 0.9.

```
In [13]: criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```

## Train Network

### Construct helper functions

In [14]:

```
def train(model, data_loader, device=torch.device('cpu')):
    """ Helper function for training one epoch of neural network.

    Args:
        model: the neural network to be trained
        data_loader: for loading the netowrk input and targets from the training dataset
        device: device used when training, default is CPU

    Returns:
        model: the trained model
        train_loss: average loss value on the entire training dataset
        train_accuracy: average accuracy on the entire training dataset
    """
    running_loss = 0.0
    correct = 0
    total = 0

    for i, data in enumerate(data_loader, 0):
        inputs, labels = data[0].to(device), data[1].to(device)
        optimizer.zero_grad()

        outputs = model(inputs.float())
        _, predicted = torch.max(outputs, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()

    train_loss = running_loss / total
    train_accuracy = correct * 100 / total

    return model, train_loss, train_accuracy
```

In [15]:

```
def test(model, data_loader, device=torch.device('cpu')):
    """ Helper function for evaluating a trained neural network on a testing set.

    Args:
        model: trained neural network
        data_loader: for loading the netowrk input and targets from the testing or validation dataset
        device: device used when training, default is CPU

    Returns:
        test_loss: average loss value on the entire testing dataset
        test_accuracy: percentage of correctly classified samples in the testing or validation dataset
    """
    running_loss = 0.0
    correct = 0
    total = 0

    with torch.no_grad():
        for i, data in enumerate(data_loader, 0):
            inputs, labels = data[0].to(device), data[1].to(device)
            outputs = model(inputs.float())

            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

            loss = criterion(outputs, labels)
            running_loss += loss.item()

    test_loss = running_loss / total
    test_accuracy = correct * 100 / total

    return test_loss, test_accuracy
```

In [16]:

```
def train_model(model, train_loader=None, valid_loader=None, n_epochs=1, stop_thr=1e-4):
```

```

""" Helper function for training a model.

Args:
    model: the neural network to be trained
    train_loader: for loading the netowrk input and targets from the training dataset
    valid_loader: for loading the netowrk input and targets from the validation dataset
    n_epochs: maximum number of epoch for training the model
    stop_thre: if the validation loss from one epoch to the next is less than this
               value, stop training

Returns:
    model: the trained model
    loss: a list of losses after each epoch of training
    acc: a list of accuracies after each epoch of training
"""
loss = dict({'train': [], 'valid': []})
acc = dict({'train': [], 'valid': []})
test_loss_pre = 100000
begin = time.time()
for epoch in range(n_epochs):
    # train
    model, train_loss, train_accuracy = train(model, train_loader)
    loss['train'].append(train_loss)
    acc['train'].append(train_accuracy)

    # validate
    test_loss, test_accuracy = test(model, valid_loader)
    loss['valid'].append(test_loss)
    acc['valid'].append(test_accuracy)

    if test_loss_pre - test_loss < stop_thr: break
    test_loss_pre = test_loss

    print("epoch ", epoch, " --- loss = ", train_loss, " acc = ", train_accuracy, \
          "time = ", time.time()-begin)

return model, loss, acc

```

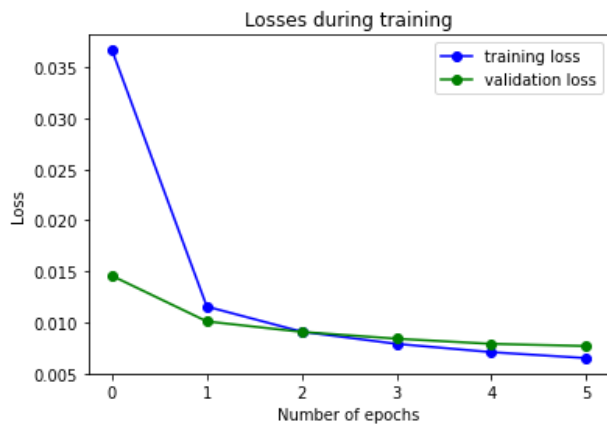
## Start Traning

```
In [17]: net, loss, acc = train_model(net, train_loader=trainloader, valid_loader=validloader, n_epochs=10)
```

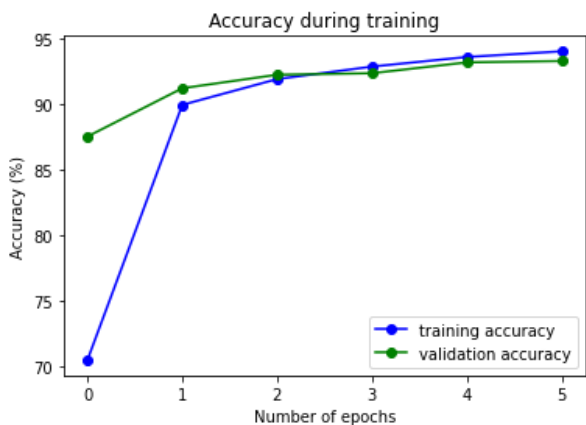
/home/laiy/.local/lib/python3.8/site-packages/torch/autograd/\_\_init\_\_.py:130: UserWarning: CUDA initialization: Found no NVIDIA driver on your system. Please check that you have an NVIDIA GPU and installed a driver from <http://www.nvidia.com/Download/index.aspx> (Triggered internally at /pytorch/c10/cuda/CUDAFunct ions.cpp:100.)

```
Variable._execution_engine.run_backward(
epoch 0 --- loss = 0.03666364541017594 acc = 70.45432692307692 time = 287.632164478302
epoch 1 --- loss = 0.011543384781032682 acc = 89.90464743589743 time = 545.9902911186218
epoch 2 --- loss = 0.009081629322549233 acc = 91.8701923076923 time = 823.1550698280334
epoch 3 --- loss = 0.007889797226749718 acc = 92.82291666666667 time = 1106.3504991531372
epoch 4 --- loss = 0.007077320216425832 acc = 93.55528846153847 time = 1359.9804983139038
epoch 5 --- loss = 0.006493620869650607 acc = 93.98717948717949 time = 1604.0112090110779
```

```
In [18]: n = range(len(loss['train']))
plt.plot(n, loss['train'], 'bo-', label="training loss")
plt.plot(n, loss['valid'], 'go-', label="validation loss")
plt.title("Losses during training")
plt.xlabel('Number of epochs')
plt.ylabel("Loss")
plt.legend(loc="upper right")
plt.savefig("training_loss.jpg")
plt.show()
```



```
In [19]: plt.plot(n, acc['train'], 'bo-', label="training accuracy")
plt.plot(n, acc['valid'], 'go-', label="validation accuracy")
plt.title("Accuracy during training")
plt.xlabel('Number of epochs')
plt.ylabel('Accuracy (%)')
plt.legend(loc="lower right")
plt.savefig("training_acc.jpg")
plt.show()
```



## Save Network

```
In [20]: PATH = './model_letters.pth'
torch.save(net.state_dict(), PATH)
```

```
In [21]: # check if model is successfully saved
net = Net()
net.load_state_dict(torch.load(PATH))
```

Out[21]: <All keys matched successfully>


## Test Network

### Predict on one batch

```
In [22]: dataiter = iter(testloader)
images, labels = dataiter.next()

print("Gound truth")
display_batch(images, labels)
```


Gound truth



```
In [23]: outputs = net(images.float())
_, predicted = torch.max(outputs, 1)

print("Prediction")
display_batch(images, predicted)
```

Prediction



## Test of testset

```
In [24]: test_loss, test_accuracy = test(net, testloader)
print("Test loss = ", test_loss)
print("Overall test accuracy = ", test_accuracy)
```

```
Test loss = 0.007702564745962333
Overall test accuracy = 93.3173076923077
```

## Test accuracy on each label

```
In [25]: def test_each_label(model, data_loader, device=torch.device('cpu')):
    """ Helper function for evaluating a trained neural network on a testing set.

    Args:
        model: trained neural network
        data_loader: for loading the netowrk input and targets from the testing dataset
        device: device used when training, default is CPU

    Returns:
        test_loss: average loss value on the entire testing dataset
        test_accuracy: percentage of correctly classified samples in the testing dataset
    """
    class_correct = np.zeros(class_size)
    class_total = np.zeros(class_size)
    class_acc = np.zeros(class_size)
    with torch.no_grad():
        for data in data_loader:
```



```

images, labels = data
outputs = model(images.float())
_, predicted = torch.max(outputs, 1)
c = (predicted == labels).squeeze()
for i in range(labels.size()[0]):
    label = labels[i]
    class_correct[label] += c[i].item()
    class_total[label] += 1


for i in range(class_size):
    if class_total[i] > 0:
        class_acc[i] = 100 * class_correct[i]/class_total[i]

return class_acc

```

```
In [26]: class_acc = test_each_label(net, testloader)
```

```
In [27]: fig = plt.figure(figsize=(20,10))
ax = fig.add_axes([0,0,1,1])
ax.bar(classes,class_acc)
ax.set_title("Accuracy (%) of each class")
fig.savefig("test_acc.jpg", bbox_inches='tight')
plt.show()
```



```
In [ ]:
```